

# QF301. Homework #5.

2021-11-22

I pledge on my honor that I have not given or received any unauthorized assistance on this assignment/examination. I further pledge that I have not copied any material from a book, article, the Internet or any other source except where I have expressly cited the source.

By filling out the following fields, you are signing this pledge. No assignment will get credit without being pledged.

Name: Siddharth Iyer

CWID: 10447455

Date: 11/15/2021

## Instructions

In this assignment, you should use R markdown to answer the questions below. Simply type your R code into embedded chunks as shown above. When you have completed the assignment, knit the document into a PDF file, and upload both the .pdf and .Rmd files to Canvas.

```
CWID = 10447455 #Place here your Campus wide ID number, this will personalize  
#your results, but still maintain the reproducible nature of using seeds.  
#If you ever need to reset the seed in this assignment, use this as your seed  
#Papers that use -1 as this CWID variable will earn 0's so make sure you change  
#this value before you submit your work.  
personal = CWID %% 10000  
set.seed(personal)#You can reset the seed at any time in your code,  
#but please always set it to this seed.
```

## Question 1 (40pt)

### Question 1.1

Use the quantmod package to obtain the daily adjusted close prices 15 different stocks. You should have at least 5 years of data for all assets. You should inspect the dates for your data to make sure you are including everything appropriately. Create a data frame of the daily log returns of all stocks. Print the first 6 lines of your data frame.

### Solution:

```
library(quantmod)  
stocks = c("MMM", "ABT", "AMD", "AAP", "AFL", "GOOG", "AMZN", "AAL", "AXP", "FDX", "AAPL", "T", "BBY", "  
  
for(stock in stocks){  
  getSymbols(stock, from="2010-01-01", to="2019-12-31", src="yahoo")  
}
```

```
price_df = data.frame(MMM$MMM.Adjusted, ABT$ABT.Adjusted, AMD$AMD.Adjusted,
  AAP$AAP.Adjusted, AFL$AFL.Adjusted, GOOG$GOOG.Adjusted,
  AMZN$AMZN.Adjusted, AAL$AAL.Adjusted, AXP$AXP.Adjusted,
  FDX$FDX.Adjusted, AAPL$AAPL.Adjusted, T$T.Adjusted,
  BBY$BBY.Adjusted, BLK$BLK.Adjusted, BK$BK.Adjusted)
colnames(price_df) <- stocks

rets_df = data.frame(diff(log(MMM$MMM.Adjusted))[-1], diff(log(ABT$ABT.Adjusted))[-1], diff(log(AMD$AMD.Adjusted))[-1],
  diff(log(AAP$AAP.Adjusted))[-1], diff(log(AFL$AFL.Adjusted))[-1], diff(log(GOOG$GOOG.Adjusted))[-1],
  diff(log(AMZN$AMZN.Adjusted))[-1], diff(log(AAL$AAL.Adjusted))[-1],
  diff(log(AXP$AXP.Adjusted))[-1], diff(log(FDX$FDX.Adjusted))[-1], ... = diff(log(AAPL$AAPL.Adjusted))[-1],
  diff(log(T$T.Adjusted))[-1], diff(log(BBY$BBY.Adjusted))[-1], diff(log(BLK$BLK.Adjusted))[-1], diff(log(BK$BK.Adjusted))[-1])
colnames(rets_df) <- stocks

head(rets_df)
```

```
##           MMM           ABT           AMD           AAP           AFL
## 2010-01-05 -0.0062831590 -0.008112179  0.001030397 -0.0059610481  0.028596976
## 2010-01-06  0.0140825380  0.005537955 -0.014523077  0.0086816850  0.008746232
## 2010-01-07  0.0007169748  0.008250426 -0.010504298 -0.0002469077  0.010675486
## 2010-01-08  0.0070211566  0.005099049 -0.004232811  0.0039446398 -0.010068227
## 2010-01-11 -0.0040397530  0.005073279 -0.031235711 -0.0098913056  0.025969850
## 2010-01-12  0.0008328642 -0.002895886 -0.055101065 -0.0175482304 -0.005140018
##           GOOG           AMZN           AAL           AXP           FDX
## 2010-01-05 -0.004413395  0.005882649  0.10724545 -0.0022014818  0.012976627
## 2010-01-06 -0.025531931 -0.018281786 -0.04231398  0.0160352184 -0.008314594
## 2010-01-07 -0.023554756 -0.017159620  0.02904362  0.0160888180 -0.010913498
## 2010-01-08  0.013243041  0.026716859 -0.01926818 -0.0007149556  0.024537071
## 2010-01-11 -0.001512745 -0.024335098 -0.01964696 -0.0115084688  0.026244121
## 2010-01-12 -0.017842125 -0.022977026  0.00790504  0.0131754886 -0.007708635
##           AAPL           T           BBY           BLK           BK
## 2010-01-05  0.001727480 -0.0049108772  0.0250627941  0.004307533  0.010548796
## 2010-01-06 -0.016034380 -0.0147406243 -0.0077953392 -0.020832191 -0.015154394
## 2010-01-07 -0.001850180 -0.0112912442  0.0157713425  0.010934231  0.041389766
## 2010-01-08  0.006626417 -0.0073524933 -0.0400300142  0.007014269  0.006114086
## 2010-01-11 -0.008860611 -0.0048093338 -0.0171851858  0.016150564 -0.017421472
## 2010-01-12 -0.011440302 -0.0003705512  0.0007643574 -0.018035496  0.001721518
```

## Question 1.2

Cluster these stocks based on the log returns.

Use K-Means Clustering with at least 20 attempts at clustering. Choose the number of clusters to use. Justify your choice in 1 paragraph (or less). Print the clusters.

## Solution:

```
model2 = kmeans(t(rets_df), nstart=20, centers = 2)
model3 = kmeans(t(rets_df), nstart=20, centers = 3)
model4 = kmeans(t(rets_df), nstart=20, centers = 4)
```

```

model5 = kmeans(t(rets_df), nstart=20, centers = 5)
model6 = kmeans(t(rets_df), nstart=20, centers = 6)
model7 = kmeans(t(rets_df), nstart=20, centers = 7)
model8 = kmeans(t(rets_df), nstart=20, centers = 8)
model9 = kmeans(t(rets_df), nstart=20, centers = 9)
model10 = kmeans(t(rets_df), nstart=20, centers = 10)
model11 = kmeans(t(rets_df), nstart=20, centers = 11)
model12 = kmeans(t(rets_df), nstart=20, centers = 12)
model13 = kmeans(t(rets_df), nstart=20, centers = 13)
model14 = kmeans(t(rets_df), nstart=20, centers = 14)

```

```
print(sort(model2$cluster))
```

```

##  MMM  ABT  AAP  AFL  GOOG  AMZN  AAL  AXP  FDX  AAPL  T  BBY  BLK  BK  AMD
##    1    1    1    1    1    1    1    1    1    1    1    1    1    1    2

```

```
print(sort(model3$cluster))
```

```

##  AMD  MMM  ABT  AAP  AFL  GOOG  AMZN  AXP  FDX  AAPL  T  BBY  BLK  BK  AAL
##    1    2    2    2    2    2    2    2    2    2    2    2    2    2    3

```

```
print(sort(model4$cluster))
```

```

##  AMD  BBY  AAL  MMM  ABT  AAP  AFL  GOOG  AMZN  AXP  FDX  AAPL  T  BLK  BK
##    1    2    3    4    4    4    4    4    4    4    4    4    4    4    4

```

```
print(sort(model5$cluster))
```

```

##  BBY  AAP  AMD  MMM  ABT  AFL  GOOG  AMZN  AXP  FDX  AAPL  T  BLK  BK  AAL
##    1    2    3    4    4    4    4    4    4    4    4    4    4    4    5

```

```
print(sort(model6$cluster))
```

```

##  BBY  AAP  GOOG  AMZN  AAPL  AAL  AMD  MMM  ABT  AFL  AXP  FDX  T  BLK  BK
##    1    2    3    3    3    4    5    6    6    6    6    6    6    6    6

```

```
print(sort(model7$cluster))
```

```

##  AAP  MMM  ABT  AFL  AXP  FDX  T  BLK  BK  GOOG  AAPL  BBY  AMD  AAL  AMZN
##    1    2    2    2    2    2    2    2    2    3    3    4    5    6    7

```

```
print(sort(model8$cluster))
```

```

##  GOOG  AAPL  MMM  ABT  T  AMZN  AAL  AAP  BBY  AFL  AXP  FDX  BLK  BK  AMD
##    1    1    2    2    2    3    4    5    6    7    7    7    7    7    8

```

```
print(sort(model9$cluster))
```

```
## AMZN AAPL AAP AFL AXP FDX BLK BK BBY MMM ABT T GOOG AAL AMD
## 1 2 3 4 4 4 4 4 5 6 6 6 7 8 9
```

```
print(sort(model10$cluster))
```

```
## AAP AFL AXP BLK BK FDX MMM ABT T AMZN AAPL GOOG AAL AMD BBY
## 1 2 2 2 2 3 4 4 4 5 6 7 8 9 10
```

```
print(sort(model11$cluster))
```

```
## AFL BLK BK AAP AMD GOOG AXP AAPL FDX MMM ABT T AAL BBY AMZN
## 1 1 1 2 3 4 5 6 7 8 8 8 9 10 11
```

```
print(sort(model12$cluster))
```

```
## MMM ABT T AMD AAP BLK BK AFL BBY AXP GOOG AAL AAPL FDX AMZN
## 1 1 1 2 3 4 4 5 6 7 8 9 10 11 12
```


```
print(sort(model13$cluster))
```

```
## AAP BLK BK AMZN MMM ABT AAL AFL FDX AMD AAPL BBY T GOOG AXP
## 1 2 2 3 4 4 5 6 7 8 9 10 11 12 13
```

```
print(sort(model14$cluster))
```

```
## GOOG BK AAPL FDX BLK AAL T MMM ABT AMZN AXP AFL AMD BBY AAP
## 1 2 3 4 5 6 7 8 8 9 10 11 12 13 14
```

```
plot(c(model2$tot.withinss, model3$tot.withinss,model4$tot.withinss, model5$tot.withinss, model6$tot.wi
```



I would use 3 clusters (note: on the graph 2 represents 3 clusters)...the largest drop in tot.withinss occurs between 2 and 3 clusters. Obviously the withinss will keep falling as you add more clusters, but the largest drop occurs between 2 and 3.

### Question 1.3


Cluster these stocks based on the log returns.

Use Hierarchical Clustering with complete distance metric and print the dendrogram. Choose the number of clusters to use. Justify your choice in 1 paragraph (or less). Print the clusters.

### Solution:

```
hc.complete.rets = hclust(dist(scale(t(rets_df))), method="complete")
plot(hc.complete.rets, main="Complete Linkage", cex=.9)
```

## Complete Linkage



```
dist(scale(t(rets_df)))
hclust(*, "complete")
```

```
sort(cutree(hc.complete.rets,6))
```

```
##  MMM  ABT  AFL  AXP  FDX  AAPL  T  BLK  BK  AMD  AAP  GOOG  AMZN  AAL  BBY
##   1    1    1    1    1    1    1    1    1    2    3    4    4    5    6
```

I would use about 3 clusters... the largest drop is in the 3rd level. If you follow the children in the 3rd level, you have about 3 clusters.

### Question 1.4


Cluster these stocks based on the log returns.

Use Hierarchical Clustering with average distance metric and print the dendrogram. Choose the number of clusters to use. Justify your choice in 1 paragraph (or less). Print the clusters.

### Solution:

```
hc.average.rets = hclust(dist(scale(t(rets_df))), method="average")
plot(hc.average.rets, main="Average Linkage", cex=.9)
```

## Average Linkage



```
dist(scale(t(rets_df)))
hclust (*, "average")
```

```
sort(cutree(hc.average.rets,3))
```

```
##  MMM  ABT  AAP  AFL  GOOG  AMZN  AXP  FDX  AAPL  T  BBY  BLK  BK  AMD  AAL
##   1    1    1    1    1    1    1    1    1    1    1    1    1    2    3
```

I would choose 3 clusters. The largest drop is in the second level, which again means 3 clusters. It's almost identical to problem 1.

### Question 1.5

Cluster these stocks based on the log returns.


Use Hierarchical Clustering with single distance metric and print the dendrogram. Choose the number of clusters to use. Justify your choice in 1 paragraph (or less). Print the clusters.

### Solution:

```
hc.single.rets = hclust(dist(scale(t(rets_df))), method="single")
cutree(hc.single.rets,2)
```

```
##  MMM  ABT  AMD  AAP  AFL  GOOG  AMZN  AAL  AXP  FDX  AAPL  T  BBY  BLK  BK
##   1    1    2    1    1    1    1    1    1    1    1    1    1    1    1
```

```
plot(hc.single.rets,main="Single Linkage",cex=.9)
```



```
dist(scale(t(rets_df)))
hclust (*, "single")
```

I would use 3 clusters because the largest drop occurs in the second level alongside BBY. This means AAL, AMD, the rest of the tree are 3 clusters.

## Question 1.6

Of the clustering methods considered, which (if any) most closely matches your intuition. Explain briefly (1 paragraph) why you choose this fit.

### Solution:

All the methods are quite bad...all of them produce clusters of uneven size and often group unrelated companies together. Also, the graphs are almost identical...so I guess the single is least computationally intensive.

## Question 2 (20pt)

### Question 2.1

Use the quantmod package to obtain the daily adjusted close prices for the SPY index and 15 different stocks. You should have at least 5 years of data for all assets. You should inspect the dates for your data to make sure you are including everything appropriately. You may use the same 15 stocks as in Question 1.



Create a data frame of the lagged daily log returns (single lag) of all stocks, lagged daily log returns (single lag) of the SPY index, and th (non-lagged) direction of the SPY index. Print the first 6 lines of your data frame.

### Solution:

```
getSymbols("SPY", from="2010-01-01", to="2019-12-31", src="yahoo")
```

```
## [1] "SPY"
```

```
spy_prices = SPY$SPY.Adjusted
spy_rets = dailyReturn(spy_prices, type="log")[-1]
head(spy_rets)
```

```
##           daily.returns
## 2010-01-05  0.0026438162
## 2010-01-06  0.0007036038
## 2010-01-07  0.0042123913
## 2010-01-08  0.0033223391
## 2010-01-11  0.0013956942
## 2010-01-12 -0.0093700320
```

```
df = rets_df
df["SPY_lagged"] = spy_rets
new_r = rep(0, 16)
df = rbind(new_r, df)
df = df[-length(df),]
df["SPY"] = ((spy_rets > 0)+0)
df = df[-1,]
head(df)
```

```
##           MMM           ABT           AMD           AAP           AFL
## 2010-01-05 -0.0062831590 -0.008112179  0.001030397 -0.0059610481  0.028596976
## 2010-01-06  0.0140825380  0.005537955 -0.014523077  0.0086816850  0.008746232
## 2010-01-07  0.0007169748  0.008250426 -0.010504298 -0.0002469077  0.010675486
## 2010-01-08  0.0070211566  0.005099049 -0.004232811  0.0039446398 -0.010068227
## 2010-01-11 -0.0040397530  0.005073279 -0.031235711 -0.0098913056  0.025969850
## 2010-01-12  0.0008328642 -0.002895886 -0.055101065 -0.0175482304 -0.005140018
##           GOOG           AMZN           AAL           AXP           FDX
## 2010-01-05 -0.004413395  0.005882649  0.10724545 -0.0022014818  0.012976627
## 2010-01-06 -0.025531931 -0.018281786 -0.04231398  0.0160352184 -0.008314594
## 2010-01-07 -0.023554756 -0.017159620  0.02904362  0.0160888180 -0.010913498
## 2010-01-08  0.013243041  0.026716859 -0.01926818 -0.0007149556  0.024537071
## 2010-01-11 -0.001512745 -0.024335098 -0.01964696 -0.0115084688  0.026244121
## 2010-01-12 -0.017842125 -0.022977026  0.00790504  0.0131754886 -0.007708635
##           AAPL           T           BBY           BLK           BK
## 2010-01-05  0.001727480 -0.0049108772  0.0250627941  0.004307533  0.010548796
## 2010-01-06 -0.016034380 -0.0147406243 -0.0077953392 -0.020832191 -0.015154394
## 2010-01-07 -0.001850180 -0.0112912442  0.0157713425  0.010934231  0.041389766
## 2010-01-08  0.006626417 -0.0073524933 -0.0400300142  0.007014269  0.006114086
```

```
## 2010-01-11 -0.008860611 -0.0048093338 -0.0171851858 0.016150564 -0.017421472
## 2010-01-12 -0.011440302 -0.0003705512 0.0007643574 -0.018035496 0.001721518
##          SPY_lagged SPY
## 2010-01-05 0.0026438162 1
## 2010-01-06 0.0007036038 1
## 2010-01-07 0.0042123913 1
## 2010-01-08 0.0033223391 1
## 2010-01-11 0.0013956942 0
## 2010-01-12 -0.0093700320 1
```

## Question 2.2

Split your data into training and testing sets (80% training and 20% test).

Train a random forest classifier using 4 variables per tree and 500 trees in order to predict the direction of the SPY index. Print the summary of your classifier. Print the test accuracy and test confusion matrix.

## Solution:

```
df = as.data.frame(df)
rownames(df) <- NULL

library(randomForest)
library(caret)

N = nrow(df)
train = sample(N, 4*N/5, replace = FALSE)
rf.model = randomForest(formula = as.factor(SPY) ~ ., data= df, proximity = TRUE, importance = TRUE, su
summary(rf.model)
```

```
##          Length Class Mode
## call          9 -none- call
## type           1 -none- character
## predicted     2010 factor numeric
## err.rate      1500 -none- numeric
## confusion       6 -none- numeric
## votes         4020 matrix numeric
## oob.times      2010 -none- numeric
## classes        2 -none- character
## importance      64 -none- numeric
## importanceSD    48 -none- numeric
## localImportance 0 -none- NULL
## proximity     4040100 -none- numeric
## ntree           1 -none- numeric
## mtry            1 -none- numeric
## forest         14 -none- list
## y              2010 factor numeric
## test           0 -none- NULL
## inbag           0 -none- NULL
## terms           3 terms  call
```

```

X_test = df[-train, -length(df)]
y_test = df[-train, length(df)]
y_test = as.vector(y_test)

y_pred = as.integer(as.vector(predict(rf.model, X_test)))

cat("Accuracy: ", 1 - sum(abs(y_test - y_pred))/length(y_pred), "\n")

```

```
## Accuracy: 0.9940358
```

```
confusionMatrix(data = factor(y_pred), reference = factor(y_test))
```

```

## Confusion Matrix and Statistics
##
##           Reference
## Prediction  0    1
##           0 245    2
##           1    1 255
##
##           Accuracy : 0.994
##           95% CI : (0.9827, 0.9988)
##       No Information Rate : 0.5109
##       P-Value [Acc > NIR] : <2e-16
##
##           Kappa : 0.9881
##
##  Mcnemar's Test P-Value : 1
##
##           Sensitivity : 0.9959
##           Specificity : 0.9922
##       Pos Pred Value : 0.9919
##       Neg Pred Value : 0.9961
##           Prevalence : 0.4891
##       Detection Rate : 0.4871
##   Detection Prevalence : 0.4911
##       Balanced Accuracy : 0.9941
##
##           'Positive' Class : 0
##

```

## Question 2.3

Using the trained classifier from Question 2.2, run Mean Decrease in Impurity (MDI) analysis. Print the feature importance of all predictors.

### Solution:

```
rf.model$importance
```

##		0	1	MeanDecreaseAccuracy	MeanDecreaseGini
## MMM	0.0060015027	0.0034401092		0.0045609268	85.761739
## ABT	0.0016476099	0.0016016670		0.0016200498	16.299760
## AMD	0.0011795684	0.0009617911		0.0010584526	6.834876
## AAP	0.0006808287	0.0003151091		0.0004756834	5.216108
## AFL	0.0057285938	0.0010186796		0.0030741250	35.773238
## GOOG	0.0028906248	0.0030872266		0.0029836392	25.295696
## AMZN	0.0016427949	0.0026352152		0.0021906401	18.133784
## AAL	0.0004346437	0.0001251010		0.0002624159	4.283076
## AXP	0.0083221938	0.0037572876		0.0057404474	39.935167
## FDX	0.0022935053	0.0016288477		0.0019121904	27.718523
## AAPL	0.0009417614	0.0018921646		0.0014765685	12.101392
## T	0.0012657447	0.0014642809		0.0013743367	9.353492
## BBY	0.0006096544	0.0007598117		0.0006906401	5.293152
## BLK	0.0069033003	0.0035303157		0.0049958064	80.444682
## BK	0.0054005420	0.0039354655		0.0045728190	40.487581
## SPY_lagged	0.4650392921	0.3636478812		0.4072827880	573.797527

## Question 2.4

Interpret the MDI feature importances (GINI) computed in Question 2.3.

Comment on the most and least important predictors (or if all predictors are of equal importance). Your response should be approximately 1 paragraph.

### Solution:

The most important factor is the lagged SPY. The least important factor is AAP. This makes sense because more often than not, AAP is in a cluster all by itself, making it a significant features but less likely to change outcomes.

This suggests it's a strongly auto regressive time series.

## Question 3 (20pt)

### Question 3.1

Consider the same data and classification problem as in Question 2. Run Mean Decrease in Accuracy (MDA) analysis on a random forest classifier. Print the feature importance of each predictor.

### Solution:

The code is already above

### Question 3.2

Interpret the MDA feature importances computed in Question 3.1.

Comment on the most and least important predictors (or if all predictors are of equal importance). Does this match the MDI feature importances found in Question 2.3. Your response should be approximately 1 paragraph.

## Solution:

...again the feature important shows that SPY\_lagged is the most important predictor, BLK is second, and the weakest predictor is AAL (not AAP). So they aren't always consistent.

## Question 4 (20pt)

### Question 4.1

Consider the same data and classification problem as in Question 2. Run Principal Component Analysis (PCA) on the 16 predictors used for the classifier in Question 2 (and 3). Print the Proportion of Variance Explained (PCA) for each principal component. How many principal components are necessary to explain 80% of the variance?

## Solution:

```
pca = prcomp(df, formula = as.factor(SPY) ~ ., scale = TRUE, subset = train)
pca$rotation
```

##	PC1	PC2	PC3	PC4	PC5
## MMM	0.2743807	-0.07167570	0.15192601	-0.01992701	-0.040564259
## ABT	0.2359376	0.01016009	0.11205894	0.15051689	-0.106807810
## AMD	0.1769015	0.16185734	-0.18794080	-0.64766853	-0.199826835
## AAP	0.1523572	-0.20798860	-0.56287493	0.58970068	0.088197340
## AFL	0.2744764	-0.19700755	0.15265244	-0.01032412	-0.035558515
## GOOG	0.2365636	0.47245437	0.07468910	0.21780152	0.065161276
## AMZN	0.2130629	0.53982693	-0.04030615	0.17645507	0.121466316
## AAL	0.1903667	-0.12833767	-0.17921552	-0.24416058	0.759889374
## AXP	0.2715488	-0.12624528	0.06693964	-0.03208079	0.090683786
## FDX	0.2650273	-0.13145479	-0.01545112	-0.12027927	0.140835680
## AAPL	0.2130289	0.42757781	-0.09293170	-0.06155431	-0.056182066
## T	0.1998274	-0.20475937	0.21514283	0.13815003	-0.397496951
## BBY	0.1560201	-0.07893888	-0.67428982	-0.16593148	-0.372969402
## BLK	0.2942162	-0.15355043	0.11684305	0.00849321	0.007336543
## BK	0.2766549	-0.25497359	0.12546211	-0.05209030	0.027536028
## SPY_lagged	0.3477590	0.02084779	0.05782065	0.01831654	-0.039578610
## SPY	0.2528571	0.04659256	0.05041752	0.04905462	-0.104103715
##	PC6	PC7	PC8	PC9	PC10
## MMM	0.096683113	-0.030532584	0.024098494	0.115477250	0.343427785
## ABT	-0.150601477	0.026841798	-0.822838785	0.252338282	-0.188448553
## AMD	-0.638941316	0.060681525	0.018846627	-0.091482805	-0.037059201
## AAP	-0.462583482	0.038746648	0.147638668	0.103633723	0.071973974
## AFL	0.029243085	-0.216071901	0.262469587	0.035558520	-0.178959472
## GOOG	0.093165238	-0.057776474	0.031152545	-0.125403829	-0.132319922
## AMZN	-0.027412976	-0.034727376	0.066886766	-0.493943020	-0.100805806
## AAL	0.161064170	0.413934276	-0.156463690	-0.043478864	-0.098744344
## AXP	-0.026926637	-0.186274169	0.088680901	0.042894739	-0.423820255
## FDX	0.029530019	-0.024375204	0.105773106	-0.067906046	0.488810818
## AAPL	0.149226797	0.192910296	0.255715202	0.728149393	0.043848356
## T	0.039029276	0.766124101	0.160275694	-0.213175305	-0.141051833

```
## BBY      0.531087209 -0.097527454 -0.129420872 -0.137001068 -0.122163546
## BLK      0.035401345 -0.177566193  0.042111042 -0.002484258  0.036240803
## BK      -0.014357903 -0.262168423  0.098540515 -0.030350250 -0.200107077
## SPY_lagged 0.006224228 -0.002459865  0.009151999  0.045127280  0.006738278
## SPY      0.023943122 -0.036858366 -0.250878106 -0.192850161  0.523297291
##          PC11      PC12      PC13      PC14      PC15
## MMM      -0.321947624 -0.367177629  0.519094631  0.332383472  0.3074216124
## ABT      -0.222236449 -0.051114161 -0.148128294 -0.051541975 -0.0909063317
## AMD      -0.019797122  0.070271264  0.066712098  0.112572122  0.0113898207
## AAP      0.013054600  0.009671203  0.045866436  0.063714628  0.0229387314
## AFL      -0.032525379 -0.197246177 -0.239215536  0.443755996 -0.6272505640
## GOOG     -0.113919199  0.655350500  0.327890825  0.212339494 -0.1221662428
## AMZN     -0.108175782 -0.478502674 -0.273147400 -0.119701556  0.1344057036
## AAL      0.097061050 -0.002298173 -0.001959062  0.185912941  0.0004394102
## AXP      0.210565521 -0.181992744  0.517494509 -0.540937382 -0.1142756463
## FDX      -0.450733047  0.203195132 -0.153160426 -0.514127173 -0.2769689400
## AAPL     0.164095246 -0.063556200 -0.190438410 -0.115359596  0.0468045323
## T        -0.005919982  0.048860905 -0.011567688 -0.068773402  0.0186028632
## BBY      -0.066435399  0.007285661  0.021495723  0.025945081  0.0005552454
## BLK      -0.018654590  0.158731047 -0.236970393  0.045584151  0.4184304458
## BK       0.164726075  0.230666335 -0.276279660 -0.008000701  0.4219585738
## SPY_lagged 0.005329062 -0.028752283 -0.013975364  0.035789232 -0.0164854365
## SPY      0.710531392 -0.007566288  0.043504675  0.012630358 -0.1562898873
##          PC16      PC17
## MMM      0.153242264 -0.12616095
## ABT      0.044077387 -0.10559512
## AMD      -0.031979900 -0.04636450
## AAP      0.011437208 -0.02979564
## AFL      -0.033296937 -0.15097749
## GOOG     0.021098201 -0.08311123
## AMZN     0.044952201 -0.07357203
## AAL      -0.012869844 -0.02368736
## AXP      -0.119497836 -0.08140704
## FDX      0.092763322 -0.06471523
## AAPL     0.037772799 -0.12271116
## T        0.009609212 -0.08686450
## BBY      0.004198411 -0.02218268
## BLK      -0.752192804 -0.12470144
## BK       0.615428185 -0.07802912
## SPY_lagged -0.021132215  0.93173162
## SPY      0.016320673 -0.10269619
```

```
pca.var = pca$sdev^2
pve = pca.var/sum(pca.var)
print("PCA explained variance: ")
```

```
## [1] "PCA explained variance: "
```

```
pve
```

```
## [1] 0.455805216 0.061580382 0.055214807 0.048426507 0.045290405 0.043608142
## [7] 0.040793674 0.035831351 0.035713246 0.032212328 0.031117674 0.024948693
## [13] 0.024080946 0.023162401 0.020868475 0.017598555 0.003747196
```

8 PCA components explain slight over 80% of the variance in the data.

### Question 4.2

Interpret the PCA computed in Question 4.1. Comment on the importance of different predictors. Does this match the MDI and MDA analysis from Questions 2 and 3? Your response should be approximately 1 paragraph.

### Solution:

PCA1: SPY\_lagged is the largest contributor PCA2: AMZN is the largest contributor in PCA2 PCA8: ABT is the largest contributor in last PCA at 80% threshold

This doesn't match the MDI and MDA exactly but the PCA itself is a linear combination of factors, so makes sense that optimizing for variance will produce different results than minimizing GINI.